

```
#importing library
import pandas as pd
from pandas import read_csv, get_dummies, DataFrame, Series
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE
from sklearn import metrics
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score
from sklearn.metrics import confusion_matrix
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import classification_report
from sklearn.metrics import f1_score

#loading the dataset
data = pd.read_csv('/content/healthcare_dataset.csv')
```

```
#to know the data info
data.info()
```

```
→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 55500 entries, 0 to 55499
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   Name             55500 non-null   object  
 1   Age              55500 non-null   int64  
 2   Gender            55500 non-null   object  
 3   Blood Type       55500 non-null   object  
 4   Medical Condition 55500 non-null   object  
 5   Date of Admission 55500 non-null   object  
 6   Doctor            55500 non-null   object  
 7   Hospital           55500 non-null   object  
 8   Insurance Provider 55500 non-null   object  
 9   Billing Amount     55500 non-null   float64 
 10  Room Number       55500 non-null   int64  
 11  Admission Type    55500 non-null   object  
 12  Discharge Date     55500 non-null   object  
 13  Medication          55500 non-null   object  
 14  Test Results        55500 non-null   object  
dtypes: float64(1), int64(2), object(12)
memory usage: 6.4+ MB
```

```
#to check the null value
data.isnull().sum()
```

	0
<b>Name</b>	0
<b>Age</b>	0
<b>Gender</b>	0
<b>Blood Type</b>	0
<b>Medical Condition</b>	0
<b>Date of Admission</b>	0
<b>Doctor</b>	0
<b>Hospital</b>	0
<b>Insurance Provider</b>	0
<b>Billing Amount</b>	0
<b>Room Number</b>	0
<b>Admission Type</b>	0
<b>Discharge Date</b>	0
<b>Medication</b>	0
<b>Test Results</b>	0

**dtype:** int64

to check value counts in columns

```
data['Blood Type'].value_counts()
```



count

**Blood Type**

A-	6969
A+	6956
AB+	6947
AB-	6945
B+	6945
B-	6944
O+	6917
O-	6877

**dtype:** int64

data['Medical Condition'].value\_counts()



count

**Medical Condition**

Arthritis	9308
Diabetes	9304
Hypertension	9245
Obesity	9231
Cancer	9227
Asthma	9185

**dtype:** int64

data['Doctor'].value\_counts()

Doctor

	count
<b>Michael Smith</b>	27
<b>John Smith</b>	22
<b>Robert Smith</b>	22
<b>Michael Johnson</b>	20
<b>James Smith</b>	20
...	...
<b>Keith Ortiz</b>	1
<b>Timothy Chapman</b>	1
<b>Terri Collins</b>	1
<b>Jill Jacobson</b>	1
<b>Scott Coleman</b>	1

40341 rows × 1 columns

**dtype:** int64

```
data['Hospital'].value_counts()
```



count

**Hospital**

<b>LLC Smith</b>	44
<b>Ltd Smith</b>	39
<b>Johnson PLC</b>	38
<b>Smith Ltd</b>	37
<b>Smith PLC</b>	36
...	...
<b>Hammond Ltd</b>	1
<b>Moran Smith and Galloway,</b>	1
<b>Winters Martin, and Neal</b>	1
<b>James-Owen</b>	1
<b>Marks, and Vazquez Jenkins</b>	1

39876 rows × 1 columns

**dtype:** int64

data['Insurance Provider'].value\_counts()



count

**Insurance Provider**

<b>Cigna</b>	11249
<b>Medicare</b>	11154
<b>UnitedHealthcare</b>	11125
<b>Blue Cross</b>	11059
<b>Aetna</b>	10913

**dtype:** int64

data['Room Number'].value\_counts()



count

## Room Number

393	181
491	177
420	175
104	175
392	171
...	...
253	112
254	111
257	111
381	110
398	109

400 rows × 1 columns

**dtype:** int64

data['Admission Type'].value\_counts()



count

## Admission Type

Elective	18655
Urgent	18576
Emergency	18269

**dtype:** int64

data['Medication'].value\_counts()



count

**Medication**

<b>Lipitor</b>	11140
<b>Ibuprofen</b>	11127
<b>Aspirin</b>	11094
<b>Paracetamol</b>	11071
<b>Penicillin</b>	11068

**dtype:** int64

data['Test Results'].value\_counts()



count

**Test Results**

<b>Abnormal</b>	18627
<b>Normal</b>	18517
<b>Inconclusive</b>	18356

**dtype:** int64

to drop unwanted columns

```
data_to_drop=['Name', 'Doctor',
              'Hospital', 'Date of Admission', 'Room Number',
              'Discharge Date']
```

data\_cleaned=data.drop(data\_to\_drop, axis=1)

data\_cleaned.info()



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 55500 entries, 0 to 55499
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   Age              55500 non-null   int64  
 1   Gender            55500 non-null   object  
 2   Blood Type        55500 non-null   object  
 3   Medical Condition 55500 non-null   object  
 4   Insurance Provider 55500 non-null   object  
 5   Billing Amount    55500 non-null   float64
```

```
6    Admission Type      55500 non-null object
7    Medication          55500 non-null object
8    Test Results        55500 non-null object
dtypes: float64(1), int64(1), object(7)
memory usage: 3.8+ MB
```

converting gender column into 1 and 2 because it is categorical values to numeric values

```
data_cleaned['Gender']=data_cleaned['Gender'].map({'Male':1,'Female':0})
```

```
data_cleaned['Gender']
```

Gender	
0	1
1	1
2	0
3	0
4	0
...	...
55495	0
55496	0
55497	0
55498	1
55499	0

55500 rows × 1 columns

**dtype:** int64

converting numerical values although it 3 values like normal, abnormal and inconclusive we convert it into two normal and abnormal

```
data_cleaned['Test Results'] = data_cleaned['Test Results'].apply(lambda x: 0 if x == 'Normal' else 1)
```

```
data_cleaned['Test Results']
```

Test Results	
0	0
1	1
2	0
3	1
4	1
...	...
<b>55495</b>	1
<b>55496</b>	0
<b>55497</b>	1
<b>55498</b>	1
<b>55499</b>	1

55500 rows × 1 columns

**dtype:** int64

data\_cleaned.info()

```
→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 55500 entries, 0 to 55499
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              55500 non-null   int64  
 1   Gender            55500 non-null   int64  
 2   Blood Type        55500 non-null   object  
 3   Medical Condition 55500 non-null   object  
 4   Insurance Provider 55500 non-null   object  
 5   Billing Amount    55500 non-null   float64 
 6   Admission Type   55500 non-null   object  
 7   Medication         55500 non-null   object  
 8   Test Results       55500 non-null   int64  
dtypes: float64(1), int64(3), object(5)
memory usage: 3.8+ MB
```

taking dummies of objects

```
data_cleaned=get_dummies(data_cleaned,['Blood Type','Medical Condition',
                                         'Insurance Provider','Admission Type','Medication'])
```

```
data_cleaned['Test Results'].value_counts()
```



count

**Test Results**

1	36983
0	18517

**dtype:** int64

dropping test result into y data and taking balance columns as X data

```
X=data_cleaned.drop('Test Results',axis=1)  
Y=data_cleaned['Test Results']
```

```
Y.value_counts()
```



count

**Test Results**

1	36983
0	18517

**dtype:** int64

## ▼ decsion tree

```
from sklearn import tree#importing tree
```

```
X_scaled=StandardScaler().fit_transform(X)  
X_train, X_test, Y_train, Y_test = train_test_split(X_scaled, Y, test_size = 0.3, random_state=1)  
X_train, Y_train =SMOTE(random_state=1).fit_resample(X_train, Y_train)#smotting the data because of imbalanced data  
dec_tree= tree.DecisionTreeClassifier(criterion='entropy',max_depth=5)#giving tree depth as per requirement  
dec_tree.fit(X_train,Y_train)
```



DecisionTreeClassifier

```
DecisionTreeClassifier(criterion='entropy', max_depth=5)
```

```
Y_pred=dec_tree.predict(X_test)
```

```
Accuracy= metrics.accuracy_score(Y_test,Y_pred)
Recall= metrics.recall_score(Y_test,Y_pred)
Precision= metrics.precision_score(Y_test,Y_pred)
print('Accuracy:', Accuracy)
print('Recall:', Recall)
print('Precision:', Precision)
f1 = f1_score(Y_test, Y_pred)
print("F1-Score:", f1)
```

```
→ Accuracy: 0.6694894894894895
    Recall: 0.9991929698708751
    Precision: 0.669772194506221
    F1-Score: 0.8019720033106625
```

```
conf_mat= metrics.confusion_matrix(Y_test,Y_pred)
```

```
conf_mat
```

```
→ array([[ 4, 5494],
       [ 9, 11143]])
```

## ▼ graph for basic dt

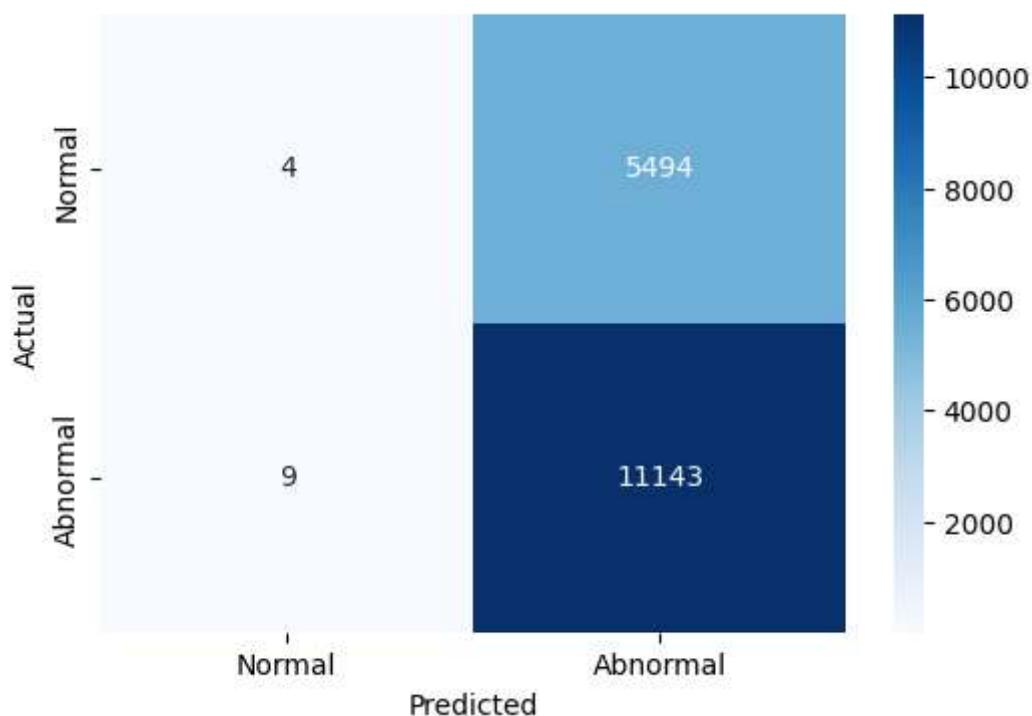
### Confusion Matrix Heatmap

```
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(6,4))
sns.heatmap(conf_mat, annot=True, fmt='d', cmap='Blues', xticklabels=["Normal", "Abnormal"],
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - Basic Decision Tree')
plt.show()
```



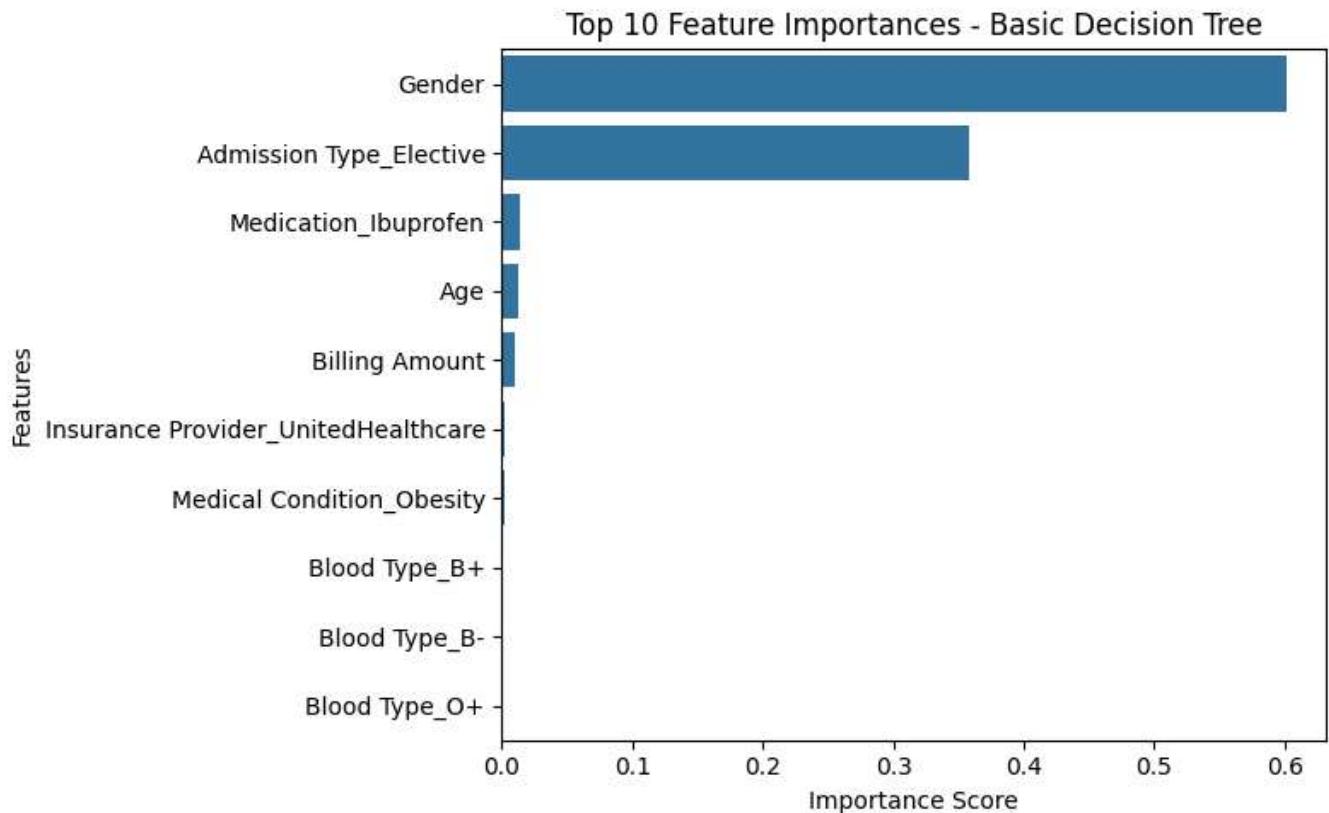
Confusion Matrix - Basic Decision Tree



## Feature Importance Bar Plot

```
feat_importance = pd.Series(dec_tree.feature_importances_, index=X.columns)
feat_importance = feat_importance.sort_values(ascending=False).head(10)
```

```
plt.figure(figsize=(8,5))
sns.barplot(x=feat_importance.values, y=feat_importance.index)
plt.title("Top 10 Feature Importances - Basic Decision Tree")
plt.xlabel("Importance Score")
plt.ylabel("Features")
plt.tight_layout()
plt.show()
```



## ▼ Tunned DT

```
from sklearn.model_selection import GridSearchCV
```

```
dec_tree_gs = tree.DecisionTreeClassifier(criterion='entropy')
depth_gs = {'max_depth':[5,10,15,20,25,30,35,40,45,46,47,48,49,50,51,52,53,54,55,60]}#to che
grid_search_gs = GridSearchCV(estimator=dec_tree_gs, param_grid=depth_gs, scoring='precision'
grid_search_gs.fit(X_train, Y_train)
best_depth_gs = grid_search_gs.best_params_
print(best_depth_gs)
```

→ { 'max\_depth': 53}

```
from sklearn.model_selection import GridSearchCV
```

```
dec_tree_gs_final = tree.DecisionTreeClassifier(criterion='entropy', max_depth=48)#got as 48
dec_tree_gs_final.fit(X_train, Y_train)
```



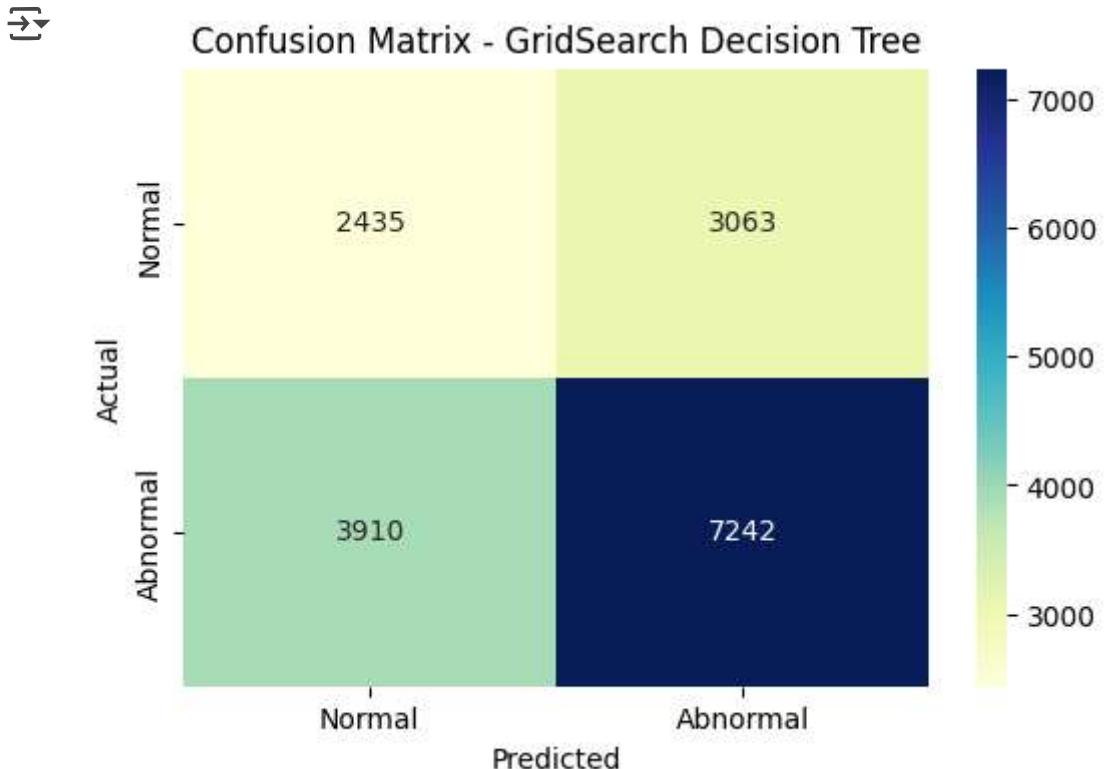
DecisionTreeClassifier  
DecisionTreeClassifier(criterion='entropy', max\_depth=48)

```
Y_pred_gs = dec_tree_gs_final.predict(X_test)
Accuracy_gs = metrics.accuracy_score(Y_test, Y_pred_gs)
Recall_gs = metrics.recall_score(Y_test, Y_pred_gs)
Precision_gs = metrics.precision_score(Y_test, Y_pred_gs)
print('Accuracy:', Accuracy_gs)
print('Recall:', Recall_gs)
print('Precision:', Precision_gs)
f1_gs = f1_score(Y_test, Y_pred_gs)
print("F1-Score:", f1_gs)
conf_mat_gs = metrics.confusion_matrix(Y_test, Y_pred_gs)
print(conf_mat_gs)
```

```
→ Accuracy: 0.5812012012012012
Recall: 0.649390243902439
Precision: 0.7027656477438137
F1-Score: 0.6750244675397307
[[2435 3063]
 [3910 7242]]
```

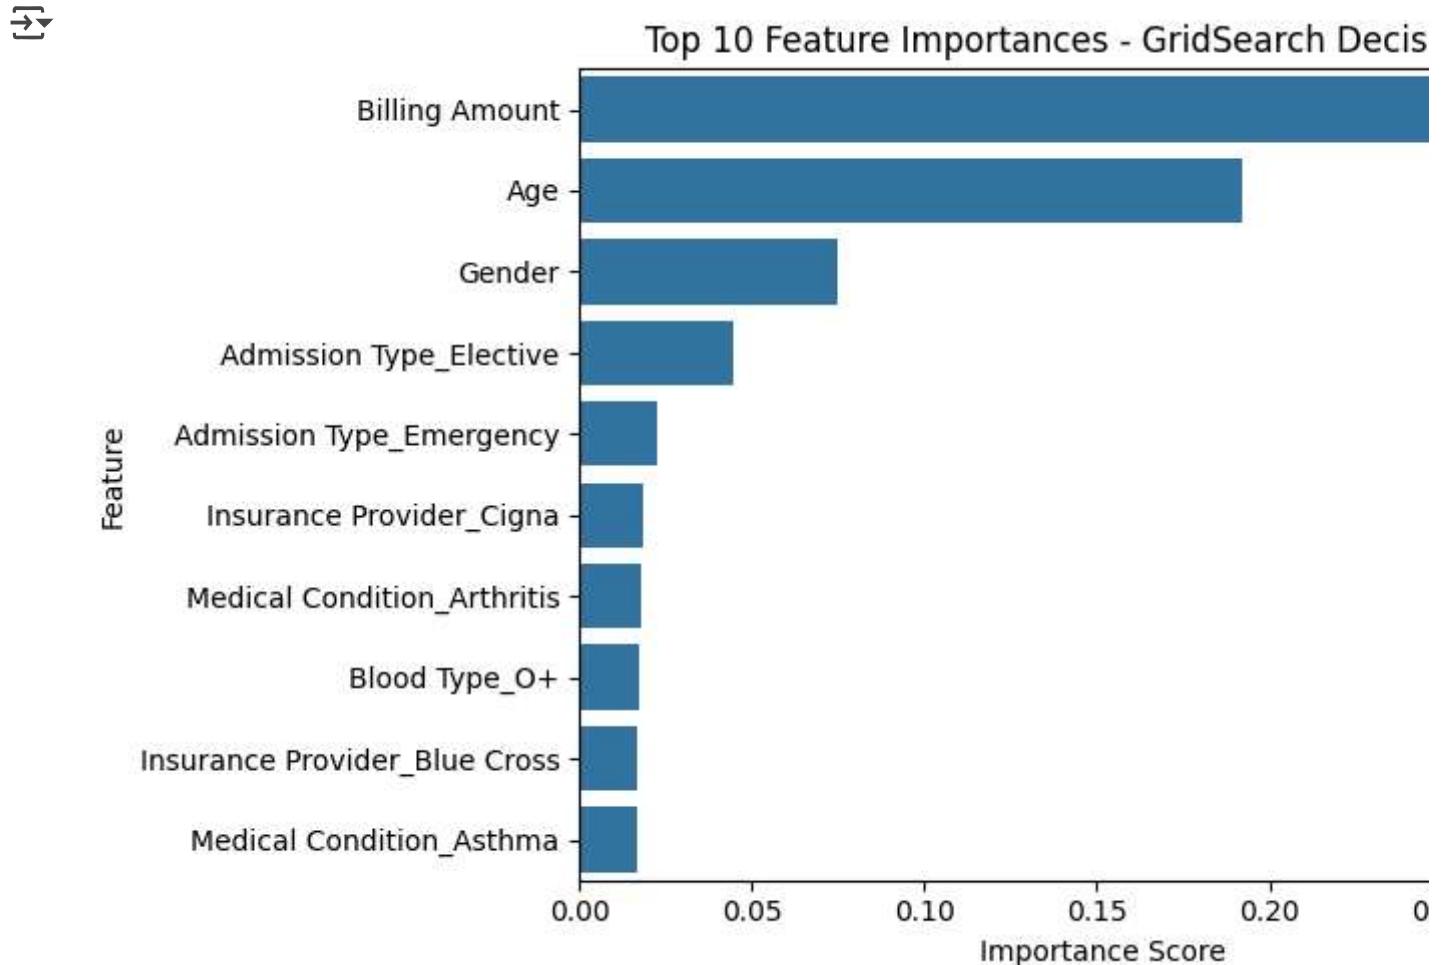
## ▼ graph for Tunned DT

```
plt.figure(figsize=(6,4))
sns.heatmap(conf_mat_gs, annot=True, fmt='d', cmap='YlGnBu', xticklabels=["Normal", "Abnormal"]
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - GridSearch Decision Tree')
plt.show()
```



```
feat_importance_gs = pd.Series(dec_tree_gs_final.feature_importances_, index=X.columns)
feat_importance_gs = feat_importance_gs.sort_values(ascending=False).head(10)

plt.figure(figsize=(8,5))
sns.barplot(x=feat_importance_gs.values, y=feat_importance_gs.index)
plt.title("Top 10 Feature Importances - GridSearch Decision Tree")
plt.xlabel("Importance Score")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```



## ▼ random forest

```
from sklearn.ensemble import RandomForestClassifier

from sklearn import ensemble
rf_model = RandomForestClassifier(n_estimators=5, criterion='entropy', max_features=None, random_state=42)
rf_model.fit(X_train, Y_train)
```

```
RandomForestClassifier
RandomForestClassifier(criterion='entropy', max_features=None, n_estimators=5,
random_state=1)
```

```
Y_pred_rf = rf_model.predict(X_test)

Accuracy_rf = accuracy_score(Y_test, Y_pred_rf)
Recall_rf = recall_score(Y_test, Y_pred_rf)
Precision_rf = precision_score(Y_test, Y_pred_rf)
f1_rf = f1_score(Y_test, Y_pred_rf)

print("Random Forest Results")
print("Accuracy:", Accuracy_rf)
print("Recall:", Recall_rf)
print("Precision:", Precision_rf)
print("F1-Score:", f1_rf)
conf_mat_rf = confusion_matrix(Y_test, Y_pred_rf)
print("Confusion Matrix:\n", conf_mat_rf)
```

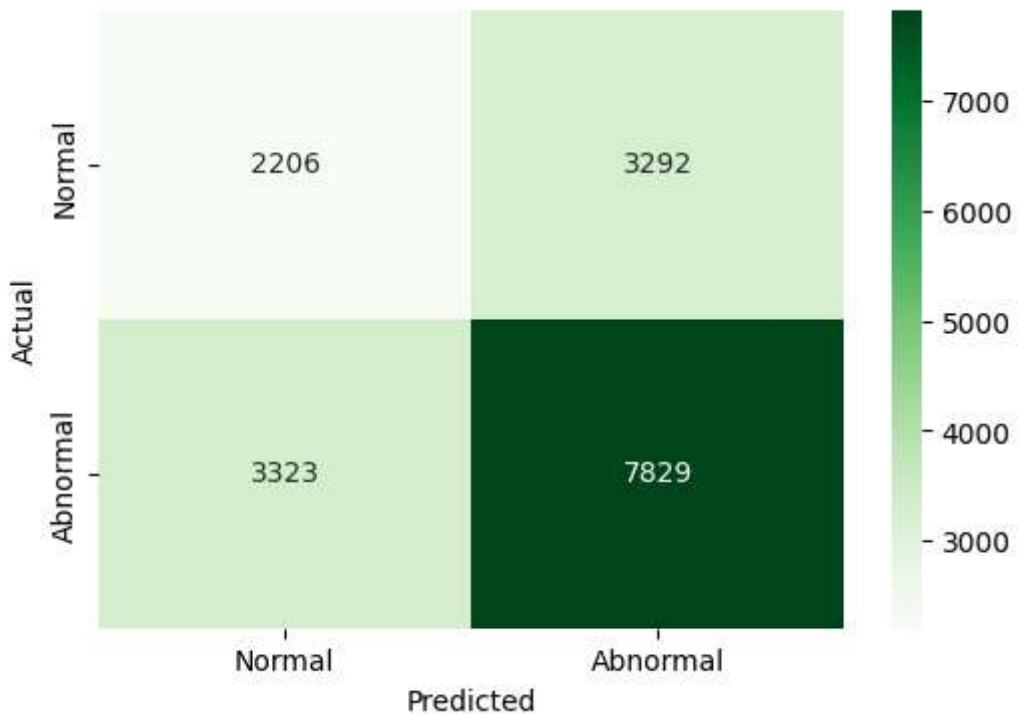
```
Random Forest Results
Accuracy: 0.6027027027027027
Recall: 0.7020265423242468
Precision: 0.7039834547252944
F1-Score: 0.703003636690163
Confusion Matrix:
[[2206 3292]
 [3323 7829]]
```

## graphs for RF

```
plt.figure(figsize=(6,4))
sns.heatmap(conf_mat_rf, annot=True, fmt='d', cmap='Greens', xticklabels=["Normal", "Abnormal"])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - Random Forest')
plt.show()
```



Confusion Matrix - Random Forest

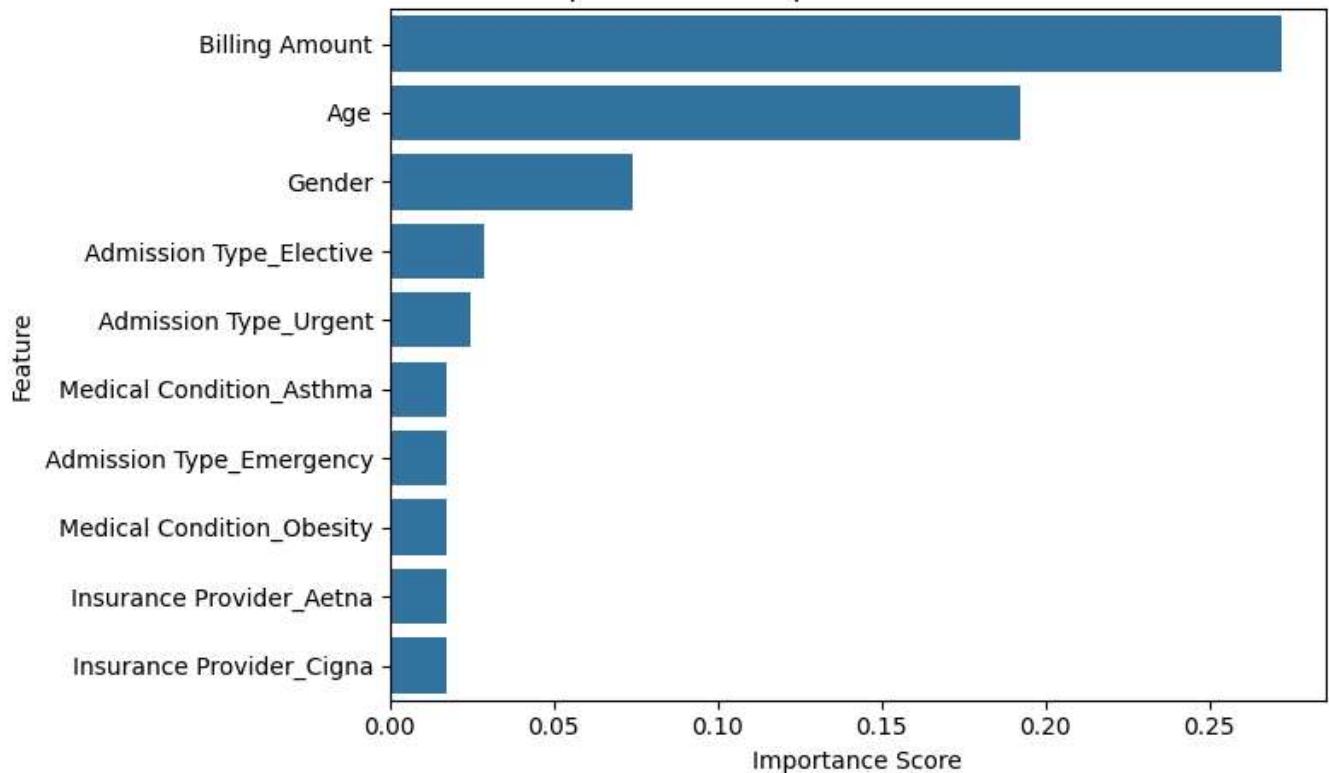


```
feat_importance_rf = pd.Series(rf_model.feature_importances_, index=X.columns)
feat_importance_rf = feat_importance_rf.sort_values(ascending=False).head(10)
```

```
plt.figure(figsize=(8,5))
sns.barplot(x=feat_importance_rf.values, y=feat_importance_rf.index)
plt.title("Top 10 Feature Importances - Random Forest")
plt.xlabel("Importance Score")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```



### Top 10 Feature Importances - Random Forest



## ▼ Tunned RF

```
rf_gs = RandomForestClassifier(criterion='entropy', max_features=None, random_state=1)
param_grid_rf = {'n_estimators': [200,300]}#to check the best estimators for that we have ch
grid_search_rf = GridSearchCV(estimator=rf_gs, param_grid=param_grid_rf, scoring='precision'
grid_search_rf.fit(X_train, Y_train)
```



### GridSearchCV

(i) (?)

**best\_estimator\_:**  
RandomForestClassifier

► RandomForestClassifier (?)

```
best_rf_params = grid_search_rf.best_params_
print("Best n_estimators:", best_rf_params)
```



Best n\_estimators: {'n\_estimators': 200}

```
rf_model_gs = RandomForestClassifier(n_estimators=best_rf_params['n_estimators'], criterion='gini', max_depth=3, min_samples_leaf=1)
# giving the best estimator as 200 but it is automatically updated from the above code
rf_model_gs.fit(X_train, Y_train)
```

```
→ RandomForestClassifier
  RandomForestClassifier(criterion='entropy', max_features=None, n_estimators=200, random_state=1)
```

```
from sklearn.metrics import f1_score
```

```
Y_pred_rf_gs = rf_model_gs.predict(X_test)
Accuracy_rf_gs = accuracy_score(Y_test, Y_pred_rf_gs)
Recall_rf_gs = recall_score(Y_test, Y_pred_rf_gs)
Precision_rf_gs = precision_score(Y_test, Y_pred_rf_gs)
f1_rf_gs = f1_score(Y_test, Y_pred_rf_gs)
```

```
print("Tuned Random Forest Results")
print("Accuracy:", Accuracy_rf_gs)
print("Recall:", Recall_rf_gs)
print("Precision:", Precision_rf_gs)
print("F1-Score:", f1_rf_gs)
conf_mat_rf_gs = confusion_matrix(Y_test, Y_pred_rf_gs)
print("Confusion Matrix:\n", conf_mat_rf_gs)
```

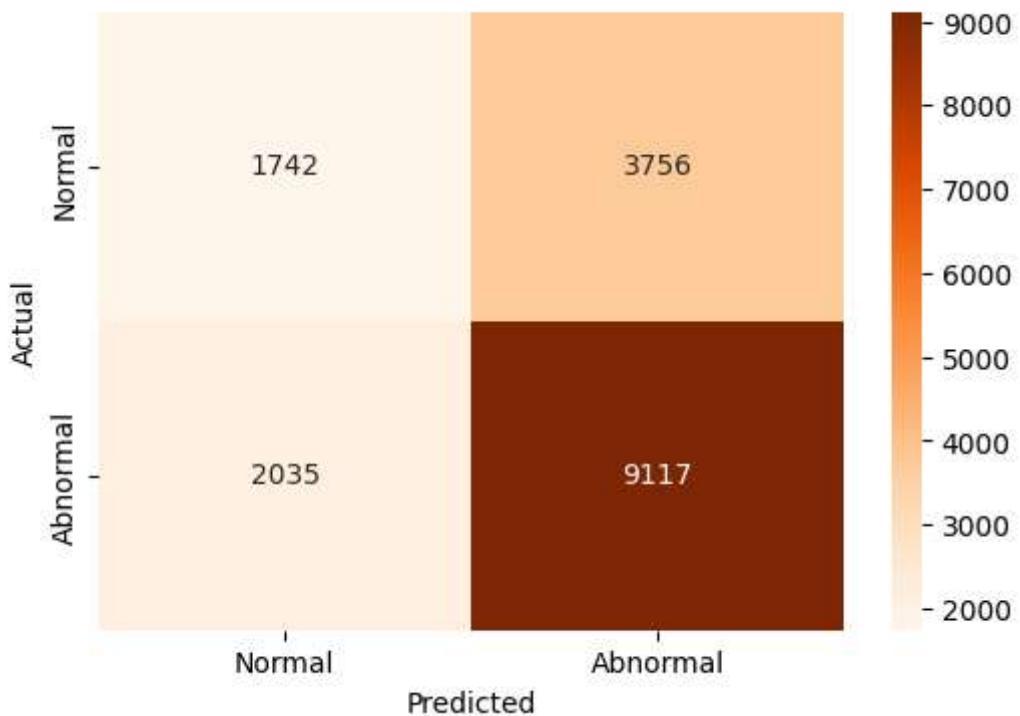
```
→ Tuned Random Forest Results
  Accuracy: 0.6521921921921922
  Recall: 0.8175215208034433
  Precision: 0.708226520624563
  F1-Score: 0.7589594172736732
  Confusion Matrix:
    [[1742 3756]
     [2035 9117]]
```

## ▼ Tunned RF graph

```
plt.figure(figsize=(6,4))
sns.heatmap(conf_mat_rf_gs, annot=True, fmt='d', cmap='Oranges', xticklabels=["Normal", "Abnormal"])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - Tuned Random Forest')
plt.show()
```



Confusion Matrix - Tuned Random Forest

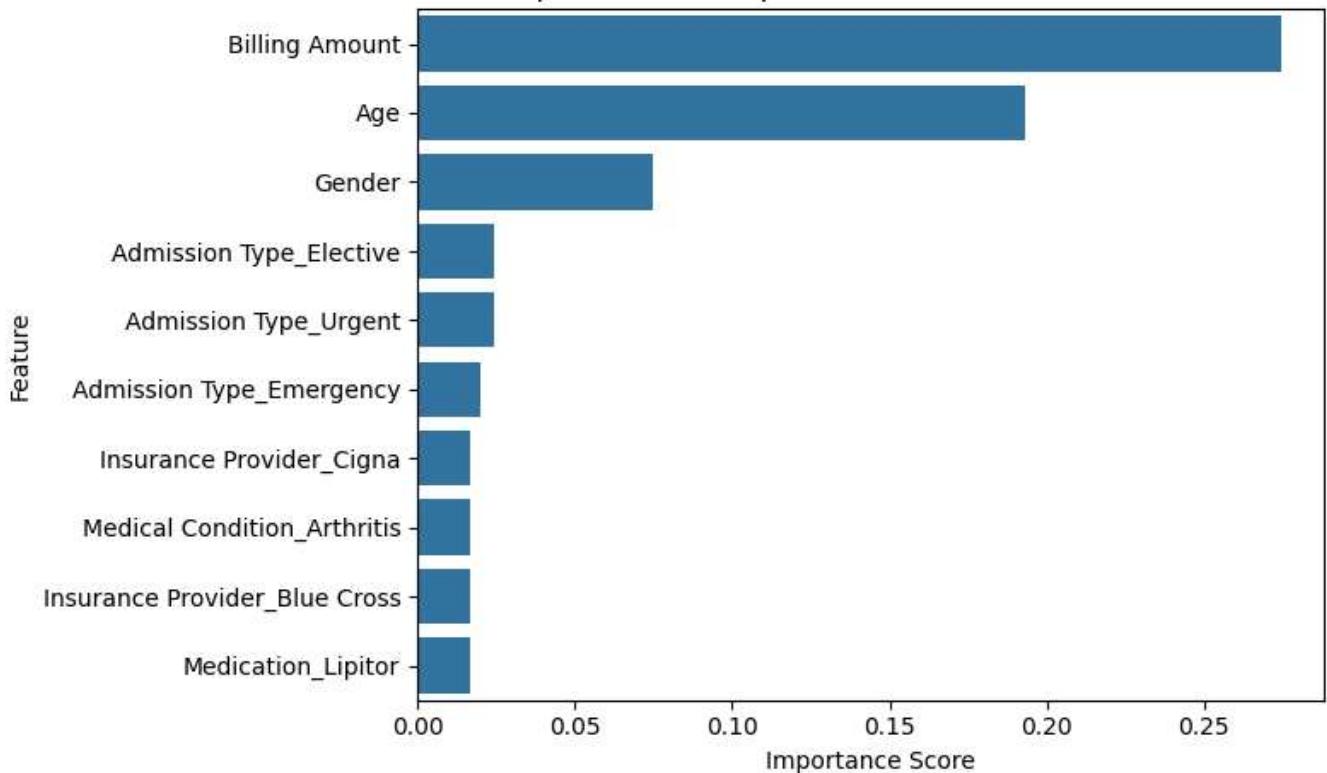


```
feat_importance_rf_gs = pd.Series(rf_model_gs.feature_importances_, index=X.columns)
feat_importance_rf_gs = feat_importance_rf_gs.sort_values(ascending=False).head(10)
```

```
plt.figure(figsize=(8,5))
sns.barplot(x=feat_importance_rf_gs.values, y=feat_importance_rf_gs.index)
plt.title("Top 10 Feature Importances - Tuned Random Forest")
plt.xlabel("Importance Score")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```



### Top 10 Feature Importances - Tuned Random Forest



## ❖ logistic Regression

```
from sklearn.linear_model import LogisticRegression

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusi
```

logreg\_model = LogisticRegression(max\_iter=1000)#giving intial value as 1000  
 logreg\_model.fit(X\_train, Y\_train)



▼ LogisticRegression i ?

LogisticRegression(max\_iter=1000)

Y\_pred\_logreg = logreg\_model.predict(X\_test)

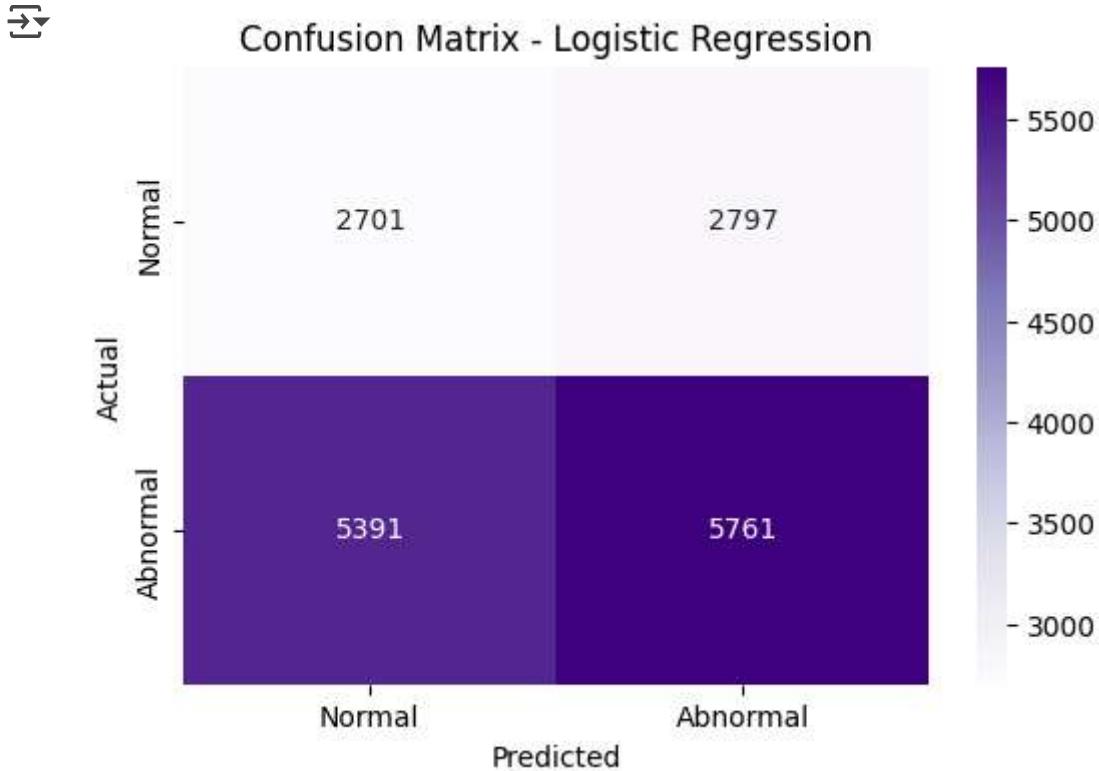
Accuracy\_logreg = accuracy\_score(Y\_test, Y\_pred\_logreg)  
 Recall\_logreg = recall\_score(Y\_test, Y\_pred\_logreg)  
 Precision\_logreg = precision\_score(Y\_test, Y\_pred\_logreg)  
 f1\_logreg = f1\_score(Y\_test, Y\_pred\_logreg)

```
print("Logistic Regression Results")
print("Accuracy:", Accuracy_logreg)
print("Recall:", Recall_logreg)
print("Precision:", Precision_logreg)
print("F1-Score:", f1_logreg)
conf_mat_logreg = confusion_matrix(Y_test, Y_pred_logreg)
print("Confusion Matrix:\n", conf_mat_logreg)
```

→ Logistic Regression Results  
Accuracy: 0.5082282282282282  
Recall: 0.5165889526542324  
Precision: 0.673171301706006  
F1-Score: 0.584576357179097  
Confusion Matrix:  
[[2701 2797]  
[5391 5761]]

## ▼ graph for logistic

```
plt.figure(figsize=(6,4))
sns.heatmap(conf_mat_logreg, annot=True, fmt='d', cmap='Purples', xticklabels=["Normal", "Abnormal"])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - Logistic Regression')
plt.show()
```



```

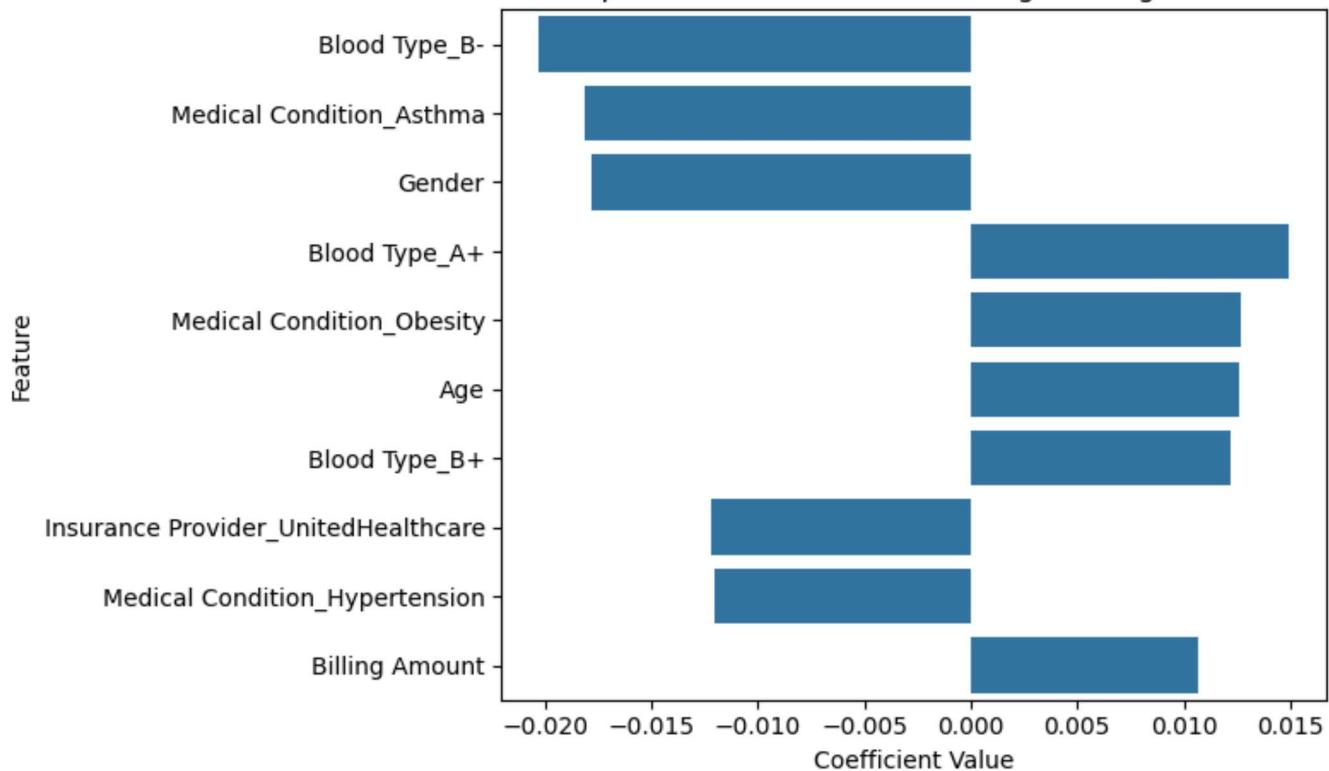
coefficients = pd.Series(logreg_model.coef_[0], index=X.columns)
coefficients = coefficients.sort_values(key=abs, ascending=False).head(10)

plt.figure(figsize=(8,5))
sns.barplot(x=coefficients.values, y=coefficients.index)
plt.title("Top 10 Feature Coefficients - Logistic Regression")
plt.xlabel("Coefficient Value")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()

```



Top 10 Feature Coefficients - Logistic Regression



## ▼ Tuned LogisticRegression

```

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

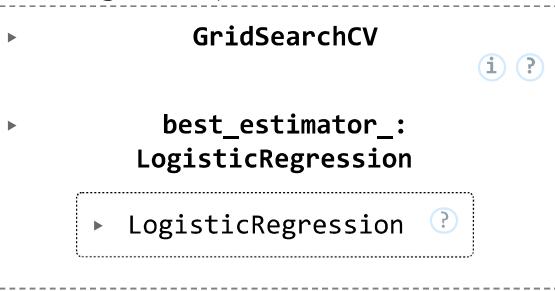
param_grid_logreg = {
    'C': [0.1, 1],
    'penalty': ['l1', 'l2']
}
#tuning the logistic for better performance

```

```
logreg_gs = LogisticRegression(max_iter=1000)
grid_search_logreg = GridSearchCV(estimator=logreg_gs, param_grid=param_grid_logreg, scoring='accuracy')
grid_search_logreg.fit(X_train, Y_train)

→ usr/local/lib/python3.11/dist-packages/sklearn/model_selection/_validation.py:528: FitFailedWarning: 0 fits failed out of a total of 20.
  The score on these train-test partitions for these parameters will be set to nan.
  If these failures are not expected, you can try to debug them by setting error_score='raise'.
  Below are more details about the failures:
-----
0 fits failed with the following error:
raceback (most recent call last):
  File "/usr/local/lib/python3.11/dist-packages/sklearn/model_selection/_validation.py", line 528, in fit
    estimator.fit(X_train, y_train, **fit_params)
  File "/usr/local/lib/python3.11/dist-packages/sklearn/base.py", line 1389, in wrapper
    return fit_method(estimator, *args, **kwargs)
           ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py", line 528, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
           ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py", line 528, in _check_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or None penalties, got l1 penalty.

  warnings.warn(some_fits_failed_message, FitFailedWarning)
/usr/local/lib/python3.11/dist-packages/sklearn/model_selection/_search.py:1108: UserWarning:
  warnings.warn(
```



```
print("Best Parameters (LogReg):", grid_search_logreg.best_params_)
```

→ Best Parameters (LogReg): {'C': 1, 'penalty': 'l2'}

```
best_logreg_model = grid_search_logreg.best_estimator_
Y_pred_logreg_gs = best_logreg_model.predict(X_test) #giving the best Best Parameters (LogReg)
```

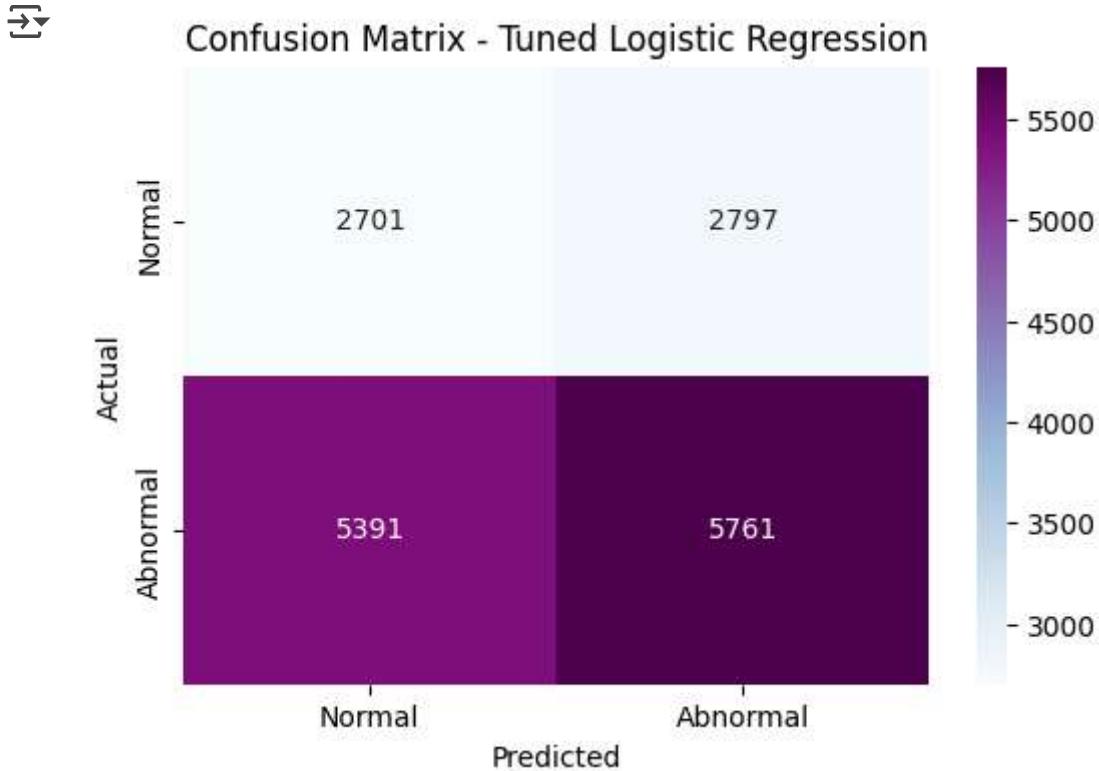
```
Accuracy_logreg_gs = accuracy_score(Y_test, Y_pred_logreg_gs)
Recall_logreg_gs = recall_score(Y_test, Y_pred_logreg_gs)
Precision_logreg_gs = precision_score(Y_test, Y_pred_logreg_gs)
f1_logreg_gs = f1_score(Y_test, Y_pred_logreg_gs)
```

```
print("GridSearch Logistic Regression Results")
print("Accuracy:", Accuracy_logreg_gs)
print("Recall:", Recall_logreg_gs)
print("Precision:", Precision_logreg_gs)
print("F1-Score:", f1_logreg_gs)
conf_mat_logreg_gs = confusion_matrix(Y_test, Y_pred_logreg_gs)
print("Confusion Matrix:\n", conf_mat_logreg_gs)
```

→ GridSearch Logistic Regression Results  
Accuracy: 0.5082282282282282  
Recall: 0.5165889526542324  
Precision: 0.673171301706006  
F1-Score: 0.584576357179097  
Confusion Matrix:  
[[2701 2797]  
[5391 5761]]

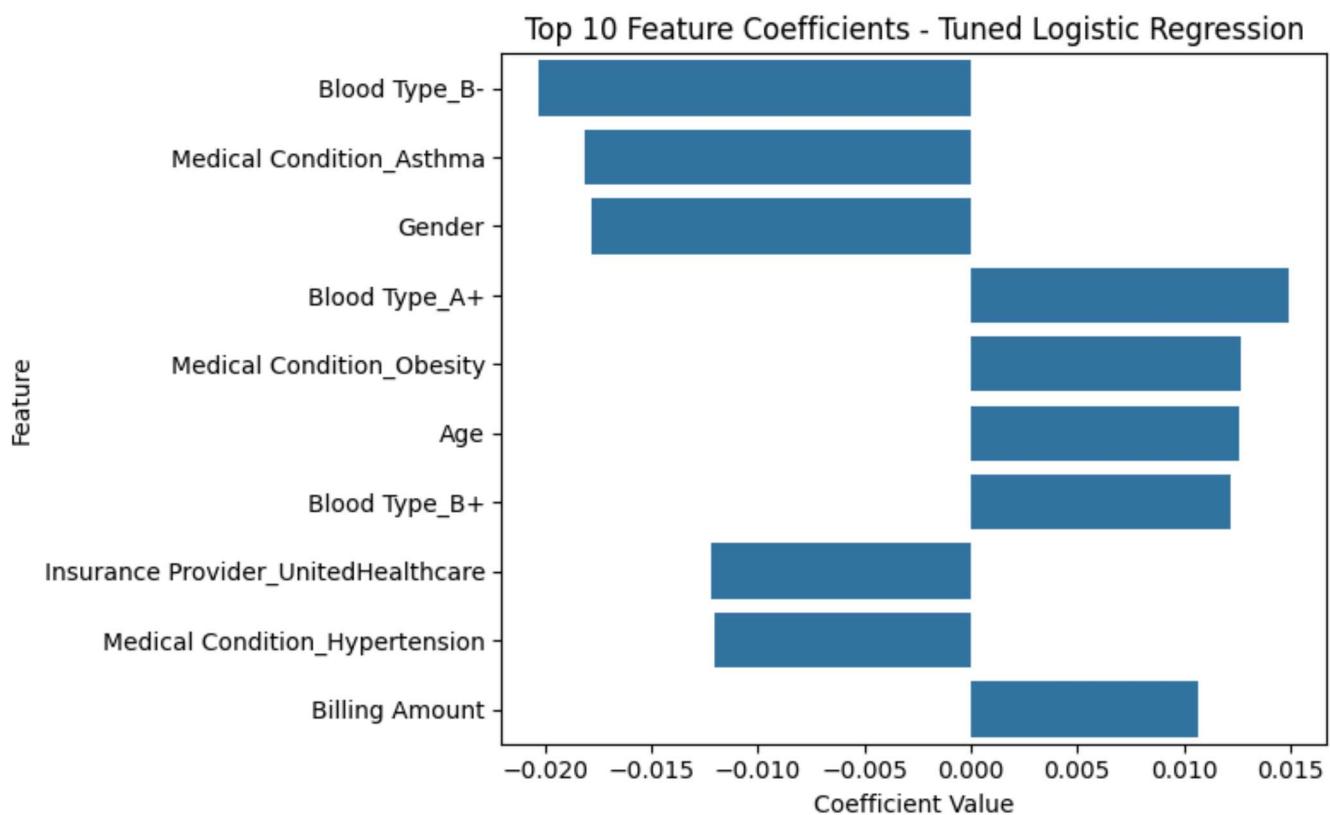
## ▼ graph for loistic

```
plt.figure(figsize=(6,4))
sns.heatmap(conf_mat_logreg_gs, annot=True, fmt='d', cmap='BuPu', xticklabels=["Normal", "Abnormal"], yticklabels=["Normal", "Abnormal"])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - Tuned Logistic Regression')
plt.show()
```



```
coefficients_logreg_gs = pd.Series(best_logreg_model.coef_[0], index=X.columns)
coefficients_logreg_gs = coefficients_logreg_gs.sort_values(key=abs, ascending=False).head(1)
```

```
plt.figure(figsize=(8,5))
sns.barplot(x=coefficients_logreg_gs.values, y=coefficients_logreg_gs.index)
plt.title("Top 10 Feature Coefficients - Tuned Logistic Regression")
plt.xlabel("Coefficient Value")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```



## ✓ final model comparison bar chart

```
import matplotlib.pyplot as plt
import numpy as np

# Define the models
models = ['Basic DT', 'Tuned DT', 'Random Forest', 'Tuned RF', 'LogReg', 'Tuned L ...]
x = np.arange(len(models))
width = 0.2

# Create lists to store the evaluation metrics for each model
accuracy_vals = [Accuracy, Accuracy_gs, Accuracy_rf, Accuracy_rf_gs, Accuracy_log ...]
precision_vals = [Precision, Precision_gs, Precision_rf, Precision_rf_gs, Precisi ...]
recall_vals = [Recall, Recall_gs, Recall_rf, Recall_rf_gs, Recall_logreg, Recall_ ...]
f1_vals = [f1, f1_gs, f1_rf, f1_rf_gs, f1_logreg, f1_logreg_gs]
```

```
# Plotting the bar chart
plt.figure(figsize=(12, 6))
plt.bar(x - 1.5*width, accuracy_vals, width, label='Accuracy', color='steelblue')
plt.title('Accuracy vs Model Type')
plt.xlabel('Model Type')
plt.ylabel('Accuracy (%)')
plt.show()
```